

Homework 6

APPM 4720/5720 Scientific Machine Learning, Fall 2024

Due date: Monday, Oct. 7 '24, before midnight, via Gradescope

Instructor: Prof. Becker
Revision date: 10/2/2024

Theme: A Toy Problem in Constrained Optimization

Instructions Collaboration with your fellow students is OK and in fact recommended, although direct copying is not allowed. The internet **is allowed for basic tasks** (e.g., looking up definitions on wikipedia, looking at documentation, looking at basic tutorials) but it is not permissible to search for solutions to the exact problem or to *post* requests for help on forums such as <http://math.stackexchange.com/>.

Setup Consider the constrained optimization problem

$$\min_{\mathbf{w}} \frac{1}{2} \|\mathbf{A}\mathbf{w} - \mathbf{y}^{(1)}\|_2^2 \quad \text{subject to} \quad \mathbf{B}\mathbf{w} = \mathbf{y}^{(2)} \quad (\text{P})$$

for $\mathbf{w} \in \mathbb{R}^d$, $\mathbf{A} \in \mathbb{R}^{n_1 \times d}$ and $\mathbf{B} \in \mathbb{R}^{n_2 \times d}$.

Problem 1: Students in 5720 only The problem (P) is simple enough that one can find a closed-form expression for it (involving matrix inverses). Derive such a closed form expression. *Hint: you could try either a change of variables, or solving the KKT equations*

Problem 2: Let $d = 20$, $n_1 = 33$ (we can allow \mathbf{A} to be over-determined) and $n_2 = 12$ (we want \mathbf{B} to be under-determined... why?), and choose \mathbf{A} and \mathbf{B} to be realizations of iid standard normal matrices. Let $\mathbf{w}^{\text{signal}}$ be the all ones vector, and define $\mathbf{y}^{(1)} = \mathbf{A}\mathbf{w}^{\text{signal}} + \sigma\mathbf{z}$ for $\sigma = 0.1$ and \mathbf{z} iid standard normal, and define $\mathbf{y}^{(2)} = \mathbf{B}\mathbf{w}^{\text{signal}}$.

Numerically find the solution to (P) using these parameters, via an optimization software package. Find a solution with as much accuracy as possible, i.e., 10^{-10} ballpark. Turn in code showing your work.

Suggestion: for Python, try the `cvxpy` package (pre-installed on colab, or install it yourself via `conda install -c conda-forge cvxpy`). This is actually an optimization framework that includes several bundled solvers. The default `SCS` solver is not super high accuracy, so I'd suggest something like `ECOS` for this problem.

Suggestion: for 5720 students, use this solver to check your answer from Problem 1.

Problem 3: Subproblems

- a) If we solve (P) via the **penalty method** (with a given penalty parameter μ), write down the subproblem, and turn in code showing how to solve the subproblem.

Note: In general one would solve the subproblem with an iterative method like gradient descent, but in this case you can solve it with linear algebra techniques.

- b) If we solve (P) via the **Augmented Lagrangian** method (with a given penalty parameter μ and dual variable ν), write down the subproblem, and turn in code showing how to solve the subproblem.

Note: Again, generally this would be iterative, but here it can be via linear algebra.

Problem 4: Since we’re using linear algebra techniques for the penalty method, there’s no need to “warm-start” the solver, hence the penalty method is as simple as solving the penalty method subproblem for a large value of μ . Plot the error as a function of μ for a reasonable range of μ values, where the error is the Euclidean norm of the difference between the **penalty method** solution $\mathbf{w}^{(\mu)}$ and the true solution from Problem 2.

Comment: If using an iterative method, then large values of μ might require many iterations. Since we’re using a linear algebra method, can you think of any downsides of using a large μ ?

Problem 5: Building on your code from part 3(b), write code that uses the **Augmented Lagrangian** method to solve (P), and turn in two plots: the first plot is the error vs μ (as in Problem 4) but using the Augmented Lagrangian method, and the second plot is the error vs the number of Augmented Lagrangian *iterations*.

Comment: You’ll want your Augmented Lagrangian code to stop iterating after it reaches some kind of tolerance criterion.

Problem 6: Students in 5720 only Make a larger problem, with $d = 1000$, $n_1 = 2000$ and $n_2 = 100$. Find the true solution to (P) using the methods from Problem 1 or 2, and then try to find a solution using both the penalty method and the Augmented Lagrangian. Can you get a solution using these methods that has error (Euclidean norm of the different \mathbf{w} solutions) less than 10^{-12} ?

Going further: Because this is a toy problem and the solution is linear in the data, one could improve the linear algebra using *iterative refinement* ideas. No need to do this for the homework.