

Physics-Informed Neural Networks (PINNs)

Tuesday, October 8, 2024

3:43 PM

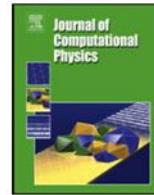
Journal of Computational Physics 378 (2019) 686–707



Contents lists available at ScienceDirect

Journal of Computational Physics

www.elsevier.com/locate/jcp



Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations



M. Raissi^a, P. Perdikaris^{b,*}, G.E. Karniadakis^a

^a Division of Applied Mathematics, Brown University, Providence, RI, 02912, USA

^b Department of Mechanical Engineering and Applied Mechanics, University of Pennsylvania, Philadelphia, PA, 19104, USA

Earlier work

- Hyuk Lee and In Seok Kang. Neural algorithm for solving differential equations. *Journal of Computational Physics*, 91(1):110–131, 1990
- Isaac E Lagaris, Aristidis Likas, and Dimitrios I Fotiadis. Artificial neural networks for solving ordinary and partial differential equations. *IEEE Trans. on neural networks*, 9(5):987–1000, 1998.

Some references:

- [https://book.sciml.ai/notes/03-Introduction to Scientific Machine Learning through Physics-Informed Neural Networks/](https://book.sciml.ai/notes/03-Introduction_to_Scientific_Machine_Learning_through_Physics-Informed_Neural_Networks/)
- Deep Learning in Computational Mechanics, Kollmannsberger et al (Springer)
- <https://www.nvidia.com/en-us/on-demand/deep-learning-for-science-and-engineering/>, co-developed by Karniadakis (PINNs and DeepOnets)
- Brittany Erickson's [MATH 607 Seminar on Physics-Informed Deep Learning](#) at U Oregon which did a lot of PINNs

I'll put more references on the github README files

Context: PDE

$$u_t + N(u, x) = 0, \quad x \in \mathbb{R}, \quad t \in [0, 1]$$

parameter

* 1D means $x \in \mathbb{R}$
but $u: \mathbb{R}^2 \rightarrow \dots$
↑ space and time

Ex: 1D Burger's Eq

$$u_t + \underbrace{\lambda_1 u u_x - \lambda_2 u_{xx}}_{N(u, x)} = 0$$

See Nick Trefethen's
"coffee table PDE" book for good, concise intro

PINNs (p. 2)

Tuesday, October 8, 2024 4:23 PM

Ex: 1D nonlinear Schrödinger Eq $\mathcal{D} = [-5, 5], t \in [0, \pi/2]$

- $i u_t + \underbrace{\frac{i}{2} u_{xx} + |u|^2 \cdot u}_{i \cdot N(u)} = 0 \quad u: \mathbb{R}^2 \rightarrow \mathbb{C}$
 $u(t, x) \in \mathbb{C}$
- initial condition: $u(t=0, x) = 2 \cdot \text{sech}(x) \rightarrow$ this is the particular I.C. in the paper
- boundary condition: $\left\{ \begin{array}{l} u(t, x=-5) = u(t, x=5) \\ u_x(t, x=-5) = u_x(t, x=5) \end{array} \right.$ (periodic B.C.)

Ex: Allen-Cahn Eq. Reaction-Diffusion $\mathcal{D} = [-1, 1], t \in [0, 1]$

$u_t - 0.0001 \cdot u_{xx} + 5u^3 - 5u = 0$ nonlinear

$u(t=0, x) = x^2 \cos(\pi x)$ I.C. from the paper

$u(t, x=-1) = u(t, x=1)$ } periodic B.C.

$u_x(t, x=-1) = u_x(t, x=1)$

Ex: Korteweg-de Vries (KdV) Eq. Burger's w, dispersive term

$u_t + \lambda_1 u u_x + \lambda_2 u_{xxx} = 0$

Ex: (toy problem) ODE not PDE

$u_t - g = 0 \quad \text{i.e.} \quad \frac{d}{dt} u(t) = g(t)$

$u(0) = u_0 \quad \text{I.C.}$

Solution via Fund. Thm. of Calculus: $u(t) = \int_0^t g(s) ds + u_0$

- Aside: in general, how do you "check your answer"?
- i.e. have the true solution?
- use an established, trusted numerical solver
 - use a problem with known analytical solution
"manufactured solution" or look up in literature

PINNs (p. 3)

Tuesday, October 8, 2024 4:38 PM

Main Ideas

① Forward problem $u_t + N(u; \lambda) = 0 \quad t \in [0, T] \quad x \in \Omega \quad + \text{I.C. + B.C.}$

assume λ parameters are known, so we won't write them explicitly
 $f := u_t + N(u) = 0$

Idea: parameterize $u: [0, T] \times \Omega \rightarrow \mathbb{R}^m$ as a neural net

Same as "Implicit Neural Representation" (INR) of SIRENS and others

then fit neural net parameters so that:

1) $f(t_f^i, x_f^i) = 0$ for a "batch" of

collocation points $\{t_f^i, x_f^i\}_{i=1}^{N_f}$

2) $u(t_u^i, x_u^i) = u^i$ for a "batch" of

points at $t=0$ (initial conditions)

and at $x \in \partial\Omega$ (boundary conditions)

We don't know u "a priori" here

here we know u since it's either I.C. or B.C.

then solve this in practice by minimizing

$c > 0$

$$\text{Loss} = \frac{1}{N_u} \sum_{i=1}^{N_u} (u(t_u^i, x_u^i) - u^i)^2 + \frac{c}{N_f} \sum_{i=1}^{N_f} f(t_f^i, x_f^i)^2$$

Not exactly penalty method but similar.

Turned root-finding into minimization

$$\text{find } F(x) = 0 \quad \min_x F(x)^2$$

i.e. $\text{Loss} = 0$ iff all constraints satisfied

DETAILS

- How to choose collocation points?

Ex:
Gauss-Chebyshev
(roots of polynomial)

Christoffel Function Sampling...

Area of approximation theory

- in small dimensions, could use a grid
- randomly ("Monte Carlo"), typically uniformly or more dense near boundary or interesting regions
- quasi-Monte Carlo (now in SciPy!) mixture of grid vs. M.C.
Aims to have high discrepancy
- space-filling designs, experimental design Stratified, Latin Hypercube, etc.
- sparse grids

PINNs (p. 4)

Tuesday, October 8, 2024 4:54 PM

How to optimize?

Since neural net has small dimensional input and output, typically not too many weights
-- so quasi-Newton (e.g. L-BFGS) often work.

e.g. 5 layers, 100 neurons each

How to evaluate f ? And its gradients?

ex: Burger's Eq., $f = u_t + (\lambda_1 u \cdot u_x - \lambda_2 u_{xx})$

u_t, u_x, u_{xx} via... backpropagation
(or any kind of AD)

$\begin{bmatrix} t \\ x \end{bmatrix} \rightarrow \text{layer 1, weights } W_1 \rightarrow \sigma \rightarrow \text{layer 2} \dots \rightarrow \text{output}$
We typically do AD for weights...
... but no technical reason we can't do AD w.r.t. these.

To train, take AD w.r.t. weights as usual

1a) Variant: Discrete-time solver (§3.2 in paper)

motivated by multi-stage methods like Runge-Kutta

Output at several stages, allows for implicit schemes

Very interesting / clever, a little weird, numerics not too convincing to me...

Larger points:

- They do not advocate solving PDEs this way!
- In high dim., they need too many collocation pts.
- In 1D, we already can solve PDE easily
- PINNs need not conserve energy, mass, ...
- Advantages:
 - 1) "easy" / "experimental"
 - 2) Easy to incorporate side information

i.e. Standard ML
w/ few data,
then add PDE
to help

PINNs (p. 5)

Tuesday, October 8, 2024 5:18 PM

[A note about I.C. + B.C.]

Vanilla PINN: $\text{loss} = \underbrace{\text{loss}_u}_{\text{loss}_{IC}} + \underbrace{c \cdot \text{loss}_f}_{\text{loss}_{BC}}$

$$c > 0 \quad \downarrow f(u, t) = 0$$

is the PDE

Initial + Boundary Conditions

choose $c=1$ or anything >0 ... in practice, one loss term may converge faster, and you need to tune c like in penalty method.

Alternative: hard-code I.C. and/or B.C.

(re-parameterization)

e.g. $u(t=0, x) = u_0(x)$ I.C., $u(t=-1) = g_{left}(t)$ B.C.

$$\mathcal{J}L = [-1, 1]$$

$u(t, +1) = g_{right}(t)$ by continuity

$$u_0(-1) = g_{left}(0)$$

$$u_0(1) = g_{right}(0)$$

define

$$u(t, x) = \begin{cases} 0 & \text{if } t=0 \text{ or } x=\pm 1 \\ t \cdot \cos(\pi/2 x) \cdot \tilde{u}(t, x) + u_0(x) & \text{otherwise} \end{cases}$$

parameterize by neural net

or do a similar style trick

$$+ \frac{1-x}{2} [g_{left}(t) - u_0(-1)] + \frac{1+x}{2} [g_{right}(t) - u_0(1)]$$

more annoying to program,

but easier to train

$$\text{so... } u(t=0, x) = 0 + u_0(x) + \frac{1-x}{2} [g_{left}(0) - u_0(-1)] + \frac{1+x}{2} [g_{right}(0) - u_0(1)] \\ = u_0(x) \quad \text{✓}$$

and

$$u(t, x=1) = 0 + u_0(1) + \frac{1-1}{2} [...] \\ + \frac{1+1}{2} [g_{right}(1) - u_0(1)] = g_{right}(1) \quad \text{✓}$$

... so we've

defined u so it automatically satisfies the constraints

[analogous to enforcing $u \geq 0$ by setting $u = \tilde{u}^2$]

and no need to fiddle with c



Caveat: I'm making this up, I don't have experience

Chris R's book.sciml.ai suggested the $u_0(x)$ I.C. trick, I added B.C. trick.

(2) Inverse Problems

- Examples are relatively low-dim inverse problems
- Many alternatives (within and without SciML)
- Includes system identification and (w)SINDy approaches

Setup: $U_t + N(U, \lambda) = 0$ but this time we know U !

ex: take measurements of a physical phenomenon

Ex: $U(t, x) = \# \text{ COVID infected people at time } t$
and location x

We may be interested in the equation (or really
parameters in the equation like " r_0 " in ODE SIR models)
for modeling or interventions

Goal: discover λ

Same idea as before: parameterize $U(t, x)$ as a neural net,
but twist: in addition to training weights ω ,
also train parameters λ

Since we have U , we set

$$U(t^i, x^i) = u^i \quad i=1, \dots, n \quad \begin{matrix} \text{collocation pts} \\ \uparrow \text{known! (may be noisy)} \end{matrix}$$

or,

$$\min_{\lambda, \omega} \frac{1}{n} \sum_{i=1}^n (U(t^i, x^i) - u^i)^2$$

(2a) Discrete-time RK version

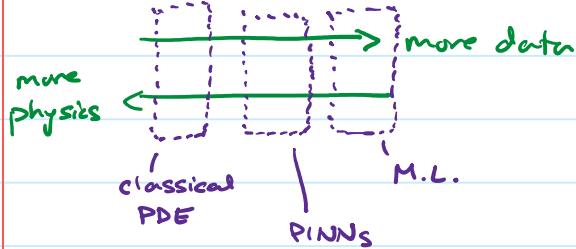
- skip -

PINNs (p. 7)

Wednesday, October 9, 2024 10:23 AM

When are PINNs useful? See Karniadakis et al., Nature Rev. Phys., '21

- If you don't need (accurate) extrapolation
- Complicated scenarios, eg. multiphysics
- Avoid complicated codes with 1000's of lines of code
- For situations where classical methods struggle
 - high-dim and/or stochastic PDE often in finance, Black-Scholes Hamilton-Jacobi-Bellman
 - For complicated geometries or to avoid meshing
 - If you have missing/gappy/low-res data on boundaries or I.C.
 - for some inverse problems where classical methods are slow due to forward simulation



"Pre-built" Software in Python see Karniadakis for some (as of '21)

eg. "DeepXDE" (w/ TensorFlow) by Karniadakis' group / NVIDIA

NeuroDiffEq is one using PyTorch