

Fourier Neural Operators (FNO)

Tuesday, October 29, 2024 1:37 PM

- Based on the 2021 ICLR paper **Fourier Neural Operator for Parametric Partial Differential Equations** by Li, Zongyi, Kovachki, Nikola, Azizzadenesheli, Kamyar, Liu, Burigede, Bhattacharya, Kaushik, Stuart, Andrew, and Anandkumar, Anima [Caltech group]
 - <https://github.com/neuraloperator/neuraloperator>, <https://zongyi-li.github.io/neural-operator/> project page
 - See their blog post <https://zongyi-li.github.io/blog/2020/fourier-pde/>
 - Extends the group's prior 2021 work **Neural Operator: Graph Kernel Network for Partial Differential Equations**
 - see their blog post <https://zongyi-li.github.io/blog/2020/graph-pde/>
 - and JMLR journal version 2023, <https://jmlr.org/papers/v24/21-1524.html>
- Neural Operator: Learning Maps Between Function Spaces
 - They have followups, like **Physics-informed neural operator for learning partial differential equations**
 - Andrew Stuart has followups on "Operator Learning Using Random Features: A Tool for Scientific Computing" (SIAM Review, 2024)
 - Anima Anandkumar has followups on "Physics-Informed Neural Operator for Learning Partial Differential Equations" (PINO) and
- For comparisons with other methods
 - Ch 6 of the book *Machine Learning in Modeling and Simulation Methods and Applications* (Timon Rabczuk and Klaus-Jürgen Bathe, eds.) "**Physics-Informed Deep Neural Operator Networks**" by Somdatta Goswami, Aniruddha Bora, Yue Yu, and George Em Karniadakis
 - compares DeepONets with FNO etc., and describes physics-informed DeepONets
 - **Operator Learning: Algorithms and Analysis**, by Nikola B. Kovachki, Samuel Lanthaler, Andrew M. Stuart (2024, <https://arxiv.org/abs/2402.15715>)
 - Math focus; describes PCA-Net, DeepONet, FNO, Random Features
 - **A mathematical guide to operator learning** by Nicolas Boullé and Alex Townsend (2023, <https://arxiv.org/abs/2312.14688>)
 - Discusses DeepONets, FNO, Deep Green networks, Graph neural op, etc.
 - Makes connection with recovering structured matrices

Idea of "operator learning"

Standard neural nets learn a mapping from X to Y

$$\hat{y} = \text{NN}(x), \quad x \in X, \quad \hat{y} \in Y$$

↑ ↑
finite dim.
vector spaces

using a mixture of :

① linear/affine layers

] "GLOBAL"

② pointwise nonlinearities

] "LOCAL"

Operator learning learns a map from

X to Y

function spaces (∞ -dim. vector spaces)

FNO (p. 2)

Tuesday, October 29, 2024 2:07 PM

First task: what are "linear layers"?

Def: Let X, Y be vector spaces. A transformation/operator is ^{"T"} $\overset{\text{(real)}}{\text{linear}}$ if $\forall f, g \in X \quad \forall \alpha \in \mathbb{R}$, then

$$T(f + \alpha g) = T(f) + \alpha T(g)$$

f a function, $h = T(f) \in Y$ also a function.

Ex $T = \text{differentiation}$, $T(f) = f'$

e.g. f is the function defined as $f(x) = 3x^2$
then $(Tf)(x) = 6x$

Ex kernel operator / integral operator

$$T(f)(x) = \int_{\Omega} K(x, y) f(y) dy$$

↑ "kernel"

e.g. $K(x, y) := \begin{cases} 1 & y \leq x \\ 0 & y > x \end{cases}$ then
 $(Tf)(x) = \int_0^x f(y) dy$
is anti-derivative

Finite-rank operators

We say a linear op. T is finite rank if $\dim(\text{range}(T)) < \infty$

- differentiation, integration are not finite rank

- an integral operator with $K(x, y) = \sum_{j=1}^r \varphi_j(x) \psi_j(y)$
Since

$$Tf \in \text{span} \{ \varphi_j \}_{j=1}^r$$

Compact operators

→ in other spaces it's more complicated

If X, Y are Hilbert space, we say a linear operator T is compact iff T is the limit (in an appropriate sense) of finite rank operators

Informally, **compact operators** are operators that we can approximate on a computer (and finite rank operators are dense among compact op.)
(banded, linear)

Also, finite rank op. are dense among a class of compact operators known as "trace class" which are equivalent to the class of integral operators (wikipedia)

⇒ Any reasonable* (**compact**) operator we wish to approximate on the computer can also be approximated by an **integral operator**.

* excludes differentiation!

So, we'll choose all linear operators in our neural operator network to be integral operators.

Motivation, part 2: GREEN'S FUNCTIONS → modern name is "Fundamental Solution"

Suppose we have a linear differential equation

$$(t) \quad \mathcal{L}(u) = f, \quad u=0 \text{ on } \partial\Omega, \quad u=u(x) \quad \text{could include time}$$

If \mathcal{L} is linear and nice enough (uniformly elliptic or parabolic; cf. Evans, §2.2.4 2010 book)

then if **G** is the **Green's Function**

which solves

$$\mathcal{L}(G(\cdot, y)) = \delta(y - \cdot) \quad] \text{ independent of } "f"$$

$$\text{then } u(x) = \int_{\Omega} G(x, y) f(y) dy$$

Why? Assuming we can commute \mathcal{L} and \int then

$$\begin{aligned} (\mathcal{L}u)(x) &= \int \mathcal{L}(G(x, y) f(y)) dy \\ &= \int \delta(y-x) f(y) dy \quad \text{convolution} \\ &= \int \delta(y) f(y+x) dy \\ &= f(x) \quad \checkmark \end{aligned}$$

... so "solution operator"

(mapping $f \mapsto u$ that solves *)

can be written as integral operators in this case.

Another reason to use integral operators

Motivation, part 3: analogy

Finite dim.

$$\vec{g} = W \cdot \vec{f}$$

$$g[] = m \begin{bmatrix} w \\ f \end{bmatrix}$$

$$g[x] = \sum_{y=1}^n w_{x,y} \cdot f[y] \quad \text{for } x \in \{1, 2, \dots, m\}$$

1D dim.

$$g = T(f) := \int_{\Omega} K(\cdot, y) f(y) dy$$

$$\text{i.e. } g(x) = \int_{\Omega} K(x, y) \cdot f(y) dy \quad \text{for } x \in \Omega$$

... going even further...

some differential equations have Green's functions of the form

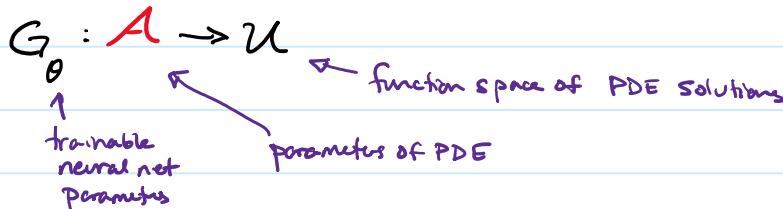
$$K(x, y) = \tilde{K}(x-y)$$

↑ usually write "K" still

... which means $\int K(x, y) f(y) dy$ is a convolution.

Back to the paper:

Goal is to learn the solution operator to a (parametric) PDE

Ex: 1D viscous Burger's Eq., $u = u(x, t)$, $x \in (0, 1)$, $t \in (0, 1]$

$$\frac{\partial}{\partial t} u + \frac{\partial}{\partial x} \left(\frac{1}{2} u^2 \right) = \nu \frac{\partial^2}{\partial x^2} u$$

 $u(x, t=0) = u_0(x)$ initial condition, u spatially periodic (boundary condition)
Goal: map $u_0 \in A$ to $u(\cdot, t=1)$ (1D function)

Ex: 2D Darcy Flow (steady-state)

$$-\nabla \cdot (\mathbf{a}(x) \nabla u(x)) = f(x) \quad x \in (0, 1)^2 \subseteq \mathbb{R}^2$$

Goal: map diffusion coefficient $a \in A$ to u (2D function)

FNO (p. 5)

Tuesday, October 29, 2024 4:15 PM

Example: Navier-Stokes 2D, viscous, incompressible in vorticity form

$$\frac{\partial}{\partial t} \omega + \mathbf{u} \cdot \nabla \omega = \nu \Delta \omega + \mathbf{f} \quad x \in (0,1)^2, \quad t \in [0, T]$$

$\underbrace{\omega}_{\text{vorticity}} := \nabla \times \mathbf{u}$ $\underbrace{\mathbf{u}}_{\text{velocity}}$

$$\nabla \cdot \mathbf{u} = 0 \quad] \text{incompressible}$$

$$\omega(x, t=0) = \omega_0(x)$$

Goal: map vorticity at time up to $t=10$ (includes ω_0 , plus more)
to vorticity up to some later time T ($T > 10$)

"Fourier layer"

FNO layer

$$\vec{V}_{t+1}(x) := \sigma \left(\underbrace{\omega \vec{V}_t(x)}_{\text{"local" (pointwise) nonlinearity}} + \underbrace{T_\theta(\vec{V}_t; \theta)}_{\substack{\text{"local"} \\ \text{integral operator}}} (x) \right)$$

\vec{V}_t $\xrightarrow{\text{layer number, not time}}$

$T_\theta(\vec{V}; \theta)(x) = \int K_\theta(x, y, \alpha(x), \alpha(y)) V(y) dy$ $\xrightarrow{\text{most general form}}$ global since $\underbrace{\text{all } y}$

learn neural net
weights θ to parameterize K_θ

$$T_\theta(\vec{V})(x) = \int K_\theta(x-y) V(y) dy \quad] \text{convolution!}$$

what they actually do

... do this via parameterizing in Fourier Space:

$$T_\theta(\vec{V})(x) = \mathcal{F}^{-1} \left(R_\theta \cdot \mathcal{F}(\vec{V}) \right) (x)$$

$\xrightarrow{\text{Inverse Fourier}}$ $\xrightarrow{\text{Fourier transform}}$

parameterize this directly

"Assume" K is periodic, so it has discrete Fourier Series,

which we then truncate to low(ish) frequencies so

it's finite.

$$|\mathbf{k}| \leq k_{\max}$$

\mathbf{k} = frequency mode

$\vec{V}(x) \in \mathbb{R}^{d_v}$, discretize input to n points, $V \in \mathbb{R}^{n \times d_v}$

Use FFT to compute $(\mathcal{F}V)[\mathbf{k}, j] \quad |\mathbf{k}| \leq k_{\max} \quad j = 1, \dots, d_v$

then $R_\theta \mathcal{F}(\vec{V})$ is

$$(R_\theta \mathcal{F}(\vec{V}))[\mathbf{k}, \lambda] = \sum_{j=1}^{d_v} R_{\mathbf{k}, \lambda, j} \cdot (\mathcal{F}V)[\mathbf{k}, j]$$

" θ " is this tensor
of size

$$(\# \text{frequencies}) \times d_v \times d_v$$

FNO (p. 6)

Tuesday, October 29, 2024 4:46 PM

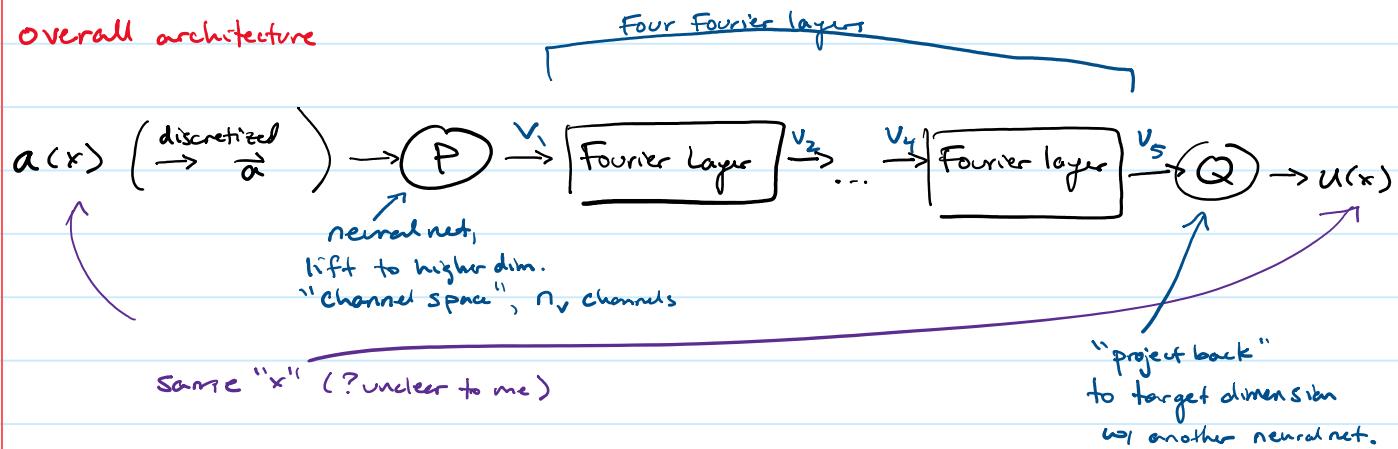
like DeepONets, while we can evaluate output u at any point x ,
our input function a (on domain D) seems to need a consistent??
discretization w/ n points,

(and simple domain + uniform discretization if we want to use FFT)

but their paper claims "discretization invariant"

since learning weights is in Fourier Space, i.e. wavenumber k_{\max} is
an absolute number, invariant to discretization.

Overall architecture



Universal Approximation

... yes, if periodic,
but size of network can grow exponentially
fast w/ ϵ unless Fourier coefficients
decay fast enough (i.e. input is smooth)

cf. Kovachki, Lanthaler,
Mishra. JMLR '21

Limitations

Assumes input, output are periodic (and same domain!?)
and uniform discretization
to use FFT

so simple domains too

Experiments

2D Navier Stokes

